



# Computer Networks

Wenzhong Li  
Nanjing University



## Chapter 5. End-to-End Protocols

- Transport Services and Mechanisms
- User Datagram Protocol (UDP)
- Transmission Control Protocol (TCP)
- TCP Congestion Control
- Real-time Transport Protocol (RTP)
- Session Initiation Protocol (SIP)
- Real Time Streaming Protocol (RTSP)



# User Datagram Protocol (UDP)

- User datagram protocol, RFC 768
- **Connectionless service** for application level processes
  - Unreliable, “best-effort” of IP
  - Each UDP segment handled independently of others
  - Delivery and duplication control not guaranteed
- **Simple and reduced overhead**
  - No connection establishment
  - No connection state at sender, receiver
  - Small segment header

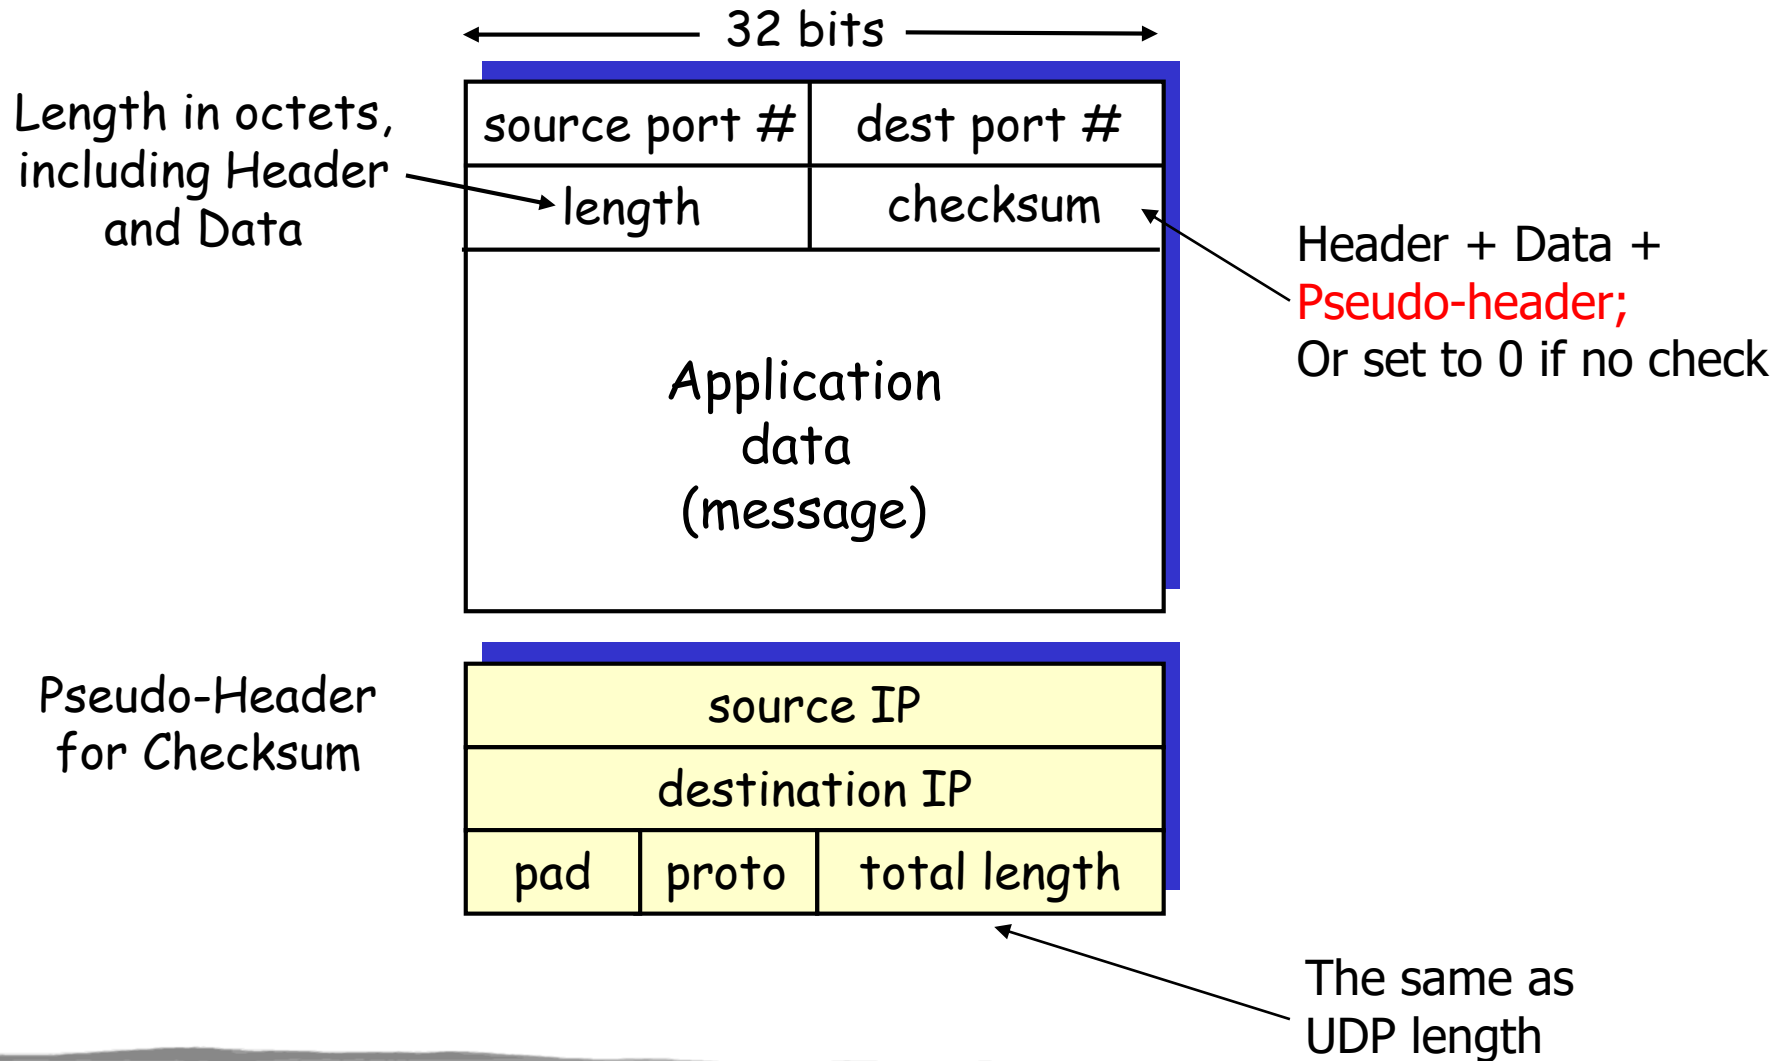


# UDP Uses

- Normal use
  - Inward data collection from sensors
  - Outward data dissemination
  - Real time applications
  - Request-Response (e.g. RPC), add reliability at application layer
- Example Apps based on UDP
  - DNS
  - SNMP



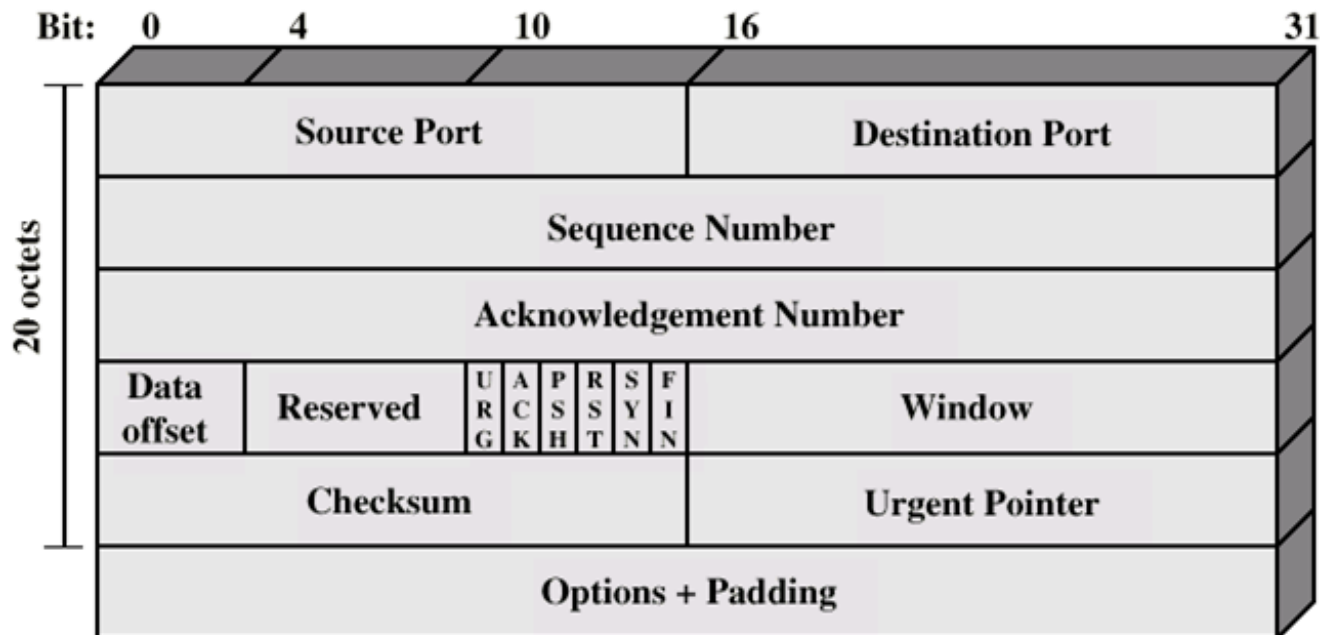
# UDP Segment Format





# Transmission Control Protocol (TCP)

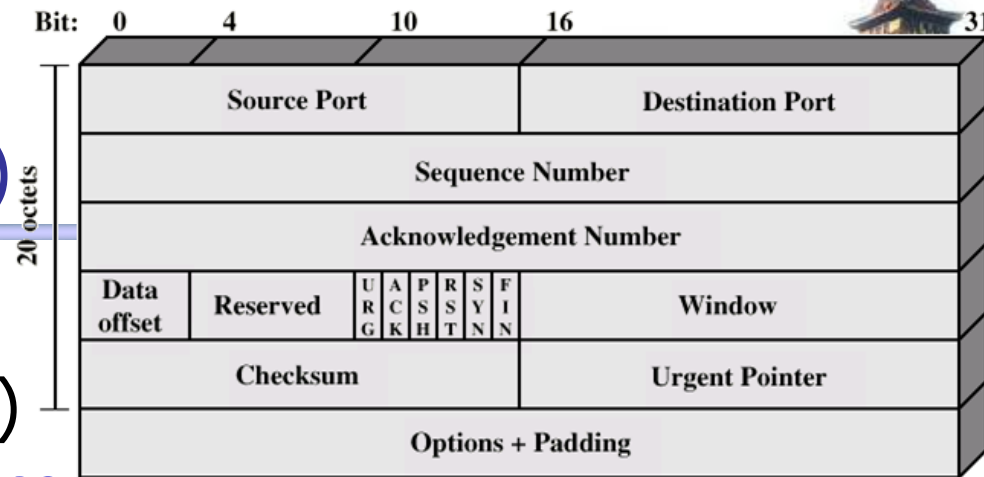
- **Reliable communication** between pairs of processes
  - Across variety of reliable and unreliable networks and internets
  - 20 bytes





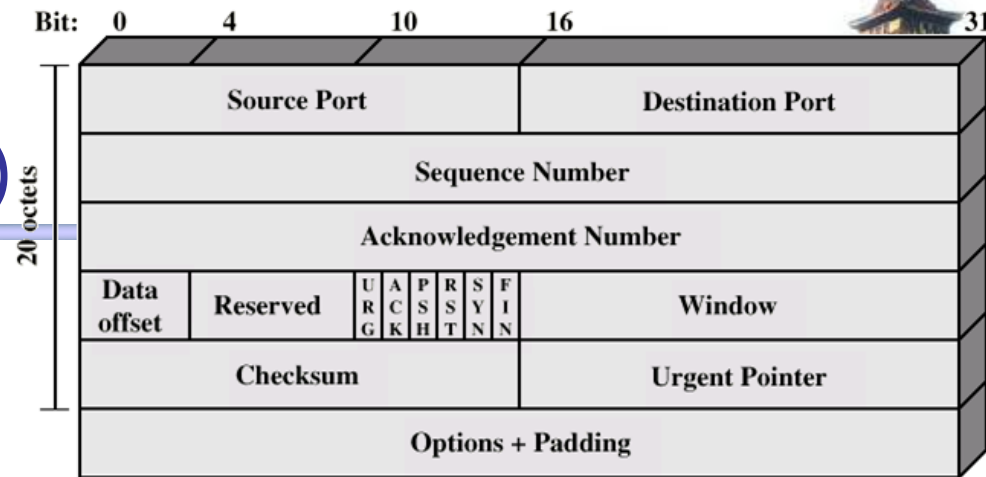
# TCP Header Fields (1)

- Source port (16 bits)
- Destination port (16 bits)
  - Identify src and dest TCP user
- Sequence number (32 bits)
  - Seq number of first data octet
  - If SYN is set, it is ISN and first data octet is ISN+1
- ACK number (32 bits)
  - Piggybacked ACK
- Window (16 bits)
  - Credit allocation in octets, i.e. rcv\_window of sender





## TCP Header Fields (2)

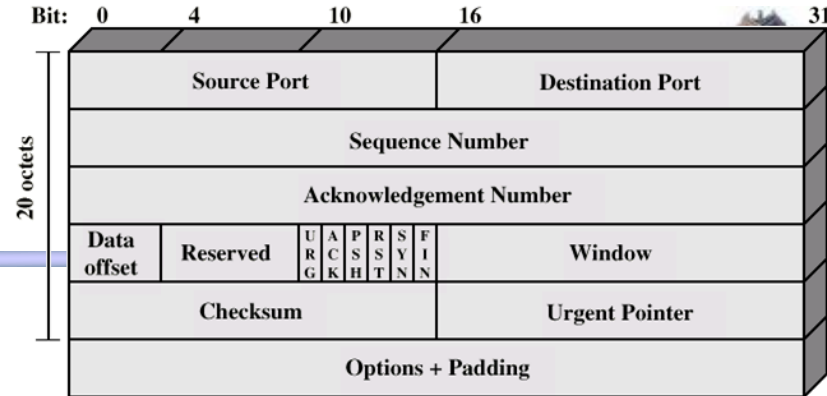


- Data offset (4 bits)
  - Number of 32-bit words in the header (报头长度)
  - Largest data offset is  $15 \times 4 = 60$  octets
- Checksum (16 bits)
  - Header + Data + Pseudo-header (src IP, dest IP, protocol No, total length)
- Reserved (6 bits)
- Options (Variable)
  - e.g. Maximum segment size the sender can accept
  - or Max value of rcv\_window





# TCP Header Fields (3)



## ■ Flags (6 bits):

- URG: urgent pointer field meaningful
- ACK: acknowledgment field meaningful
- PSH: push function
- RST: reset the connection
- SYN: synchronize the sequence number
- FIN: no more data from sender

## ■ Urgent Pointer (16 bits)

- Points to last octet in a sequence of urgent data



# Parameters Passed to IP

- TCP passes **QOS parameters** down to IP
  - Precedence
  - Normal delay / low delay
  - Normal throughput / high throughput
  - Normal reliability / high reliability
- IPv4 "Type of Service" or IPv6 "Traffic Class"



# TCP Service Request Primitives

| Primitive                    | Parameters   | Description  |
|------------------------------|--|--|
| Unspecified Passive Open     | source-port, [timeout], [timeout-action], [precedence], [security-range]   | Listen for connection attempt at specified security and precedence from any remote destination.                          |
| Fully Specified Passive Open | source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security-range]                                      | Listen for connection attempt at specified security and precedence from specified destination.                           |
| Active Open                  | source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security]  | Request connection at a particular security and precedence to a specified destination.                                   |
| Active Open with Data        | source-port, destination-port, destination-address, [timeout], [timeout-action], [precedence], [security], data, data-length, PUSH-flag, URGENT-flag | Request connection at a particular security and precedence to a specified destination and transmit data with the request |
| Send                         | local-connection-name, data, data-length, PUSH-flag, URGENT-flag, [timeout], [timeout-action]  | Transfer data across named connection  |
| Allocate                     | local-connection-name, data-length   | Issue incremental allocation for receive data to TCP   |
| Close                        | local-connection-name  | Close connection gracefully  |
| Abort                        | local-connection-name  | Close connection abruptly  |
| Status                       | local-connection-name  | Query connection status  |



# TCP Service Response Primitives

| Primitive       | Parameters  | Description  |
|-----------------|---|--|
| Open ID         | local-connection-name, source-port, destination-port*, destination-address*,  | Informs TCP user of connection name assigned to pending connection requested in an Open primitive        |
| Open Failure    | local-connection-name   | Reports failure of an Active Open request  |
| Open Success    | local-connection-name   | Reports completion of pending Open request   |
| Deliver         | local-connection-name, data, data-length, URGENT-flag   | Reports arrival of data  |
| Closing         | local-connection-name   | Reports that remote TCP user has issued a Close and that all data sent by remote user has been delivered |
| Terminate       | local-connection-name, description  | Reports that the connection has been terminated; a description of the reason for termination is provided |
| Status Response | local-connection-name, source-port, source-address, destination-port, destination-address, connection-state, receive-window, send-window, amount-awaiting-ACK, amount-awaiting-receipt, urgent-state, precedence, security, timeout | Reports current status of connection   |
| Error           | local-connection-name, description  | Reports service-request or internal error  |

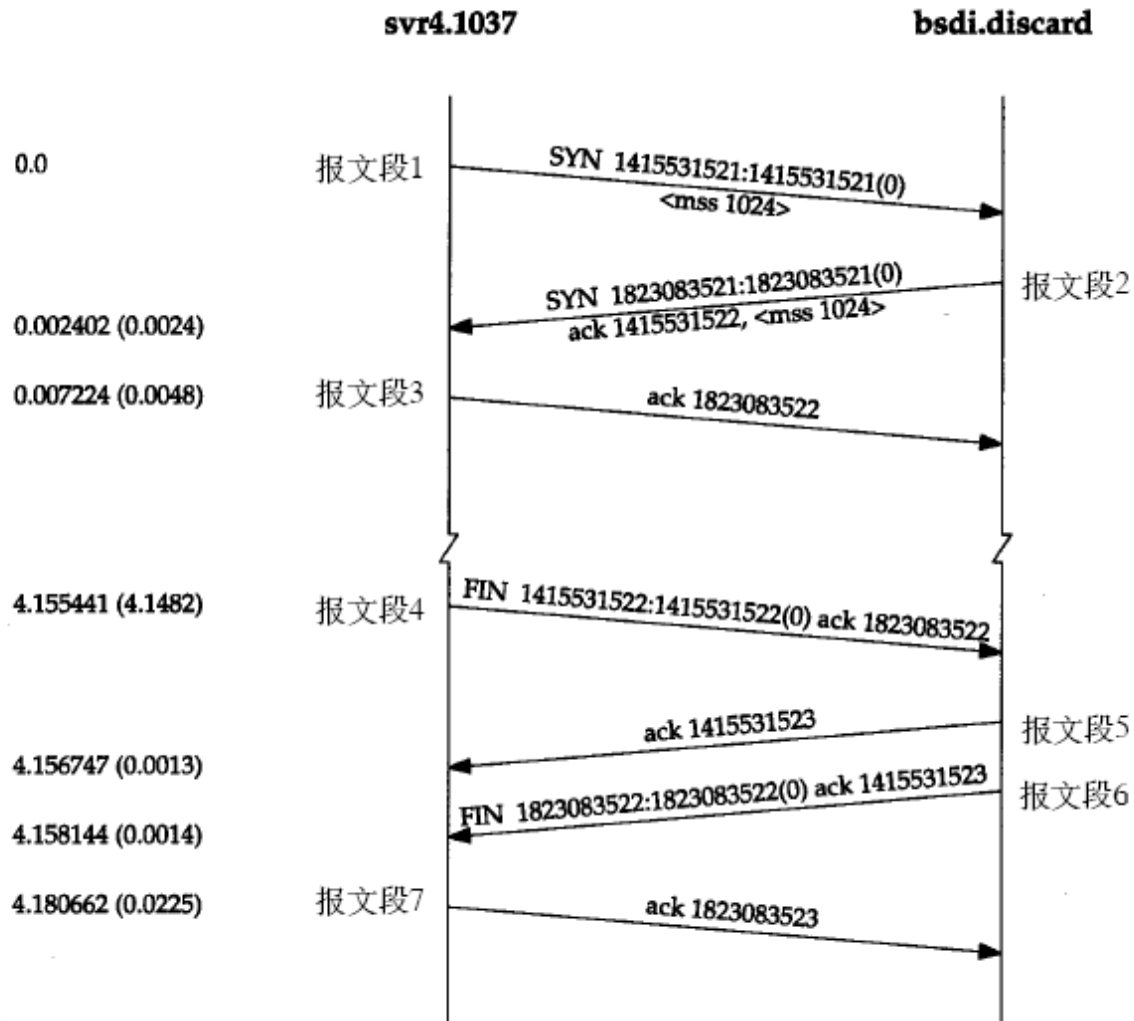


# TCP Mechanisms (1)

- Connection establishment
  - 3-way handshake
  - Between pairs of ports
  - One port can connect to multiple destination ports
- Connection termination
  - Graceful termination: `CLOSE + FIN`
  - Abrupt termination: `ABORT + RST`



# Illustration of Connection establishment & Termination





# TCP Mechanisms (2)

## ■ Data transfer

- Logical stream of octets
- Octets numbered modulo  $2^{32}$
- Flow control by credit allocation of number of octets
- Data buffered at sender and receiver
- User sets PUSH to force data transmission immediately
- User may specify a block of data as urgent



# Implementation Policy

- Send
- Deliver
- Accept
- Retransmit
- Acknowledge





# Send

- If no *PUSH* or *CLOSE*, TCP entity transmits at its own convenience
- Data issued by TCP user buffered at transmit buffer
  - May construct segment per data batch
  - May wait for certain amount of data



# Deliver

- In absence of *PUSH*, TCP entity delivers data at own convenience
- May deliver as each segment in order received
  - Deliveries (I/O interrupts) are frequent and small
- May buffer data from more than one segment
  - Deliveries are infrequent and large



# Accept

- Segments may arrive **out of order**
- In order
  - Only accept segments in order
  - Discard out of order segments
  - Makes for a simpler implementation
- In windows
  - Accept all segments **within receive window**
  - Can reduce retransmission



# Retransmit

- TCP entity maintains queue of segments **transmitted but not acknowledged**
- TCP will retransmit if not ACKed in given time
  - **First only**: one timer a queue, reset the timer after retransmission of first segment in queue
  - **Batch**: one timer a queue, reset after retransmission of all segments in queue
  - **Individual**: one timer each segment, reset after retransmission



# Fast Retransmit

- Time-out period often relatively long
- Detect lost segments via **duplicate ACKs**
  - If segment is lost, there will likely be many duplicate ACKs
- If a TCP entity receives 3 ACKs for the same data, then segments after ACKed data must be lost
  - Trigger **fast retransmit**: resend segment before timer expires



# Acknowledgement

## ■ Immediate or Cumulative Event at Receiver

## TCP Receiver action

Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed

Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK

Arrival of in-order segment with expected seq #. One other segment has ACK pending

Immediately send single cumulative ACK, ACKing both in-order segments

Arrival of out-of-order segment higher-than-expect seq. # . Gap detected

Immediately send duplicate ACK, indicating seq. # of next expected byte

Arrival of segment that partially or completely fills gap

Immediate send ACK, provided that segment sat at lower end of gap



# TCP congestion control



# TCP Congestion Control

- Congestion control
  - Too many sources sending too much data too fast for Internet to handle
  - RFC 1122, Requirements for Internet hosts
- End to end control, **no network assistance**
  - Retransmission timer
  - Window management





# Retransmission Timer Management

- Estimate round trip delay by **observing delay pattern**
  - Simple average
  - Exponential average
- **Set timer** to value somewhat greater than estimate
  - RFC 793
  - RTT Variance Estimation (Jacobson's algorithm)
- How to **set timer after retransmission**
  - Exponential RTO backoff algorithm
- **When to sample** the round trip delay
  - Karn's Algorithm



# Simple Average

## ■ Term

- $RTT(i)$ : round-trip time observed for the  $i^{\text{th}}$  transmitted segment
- $ARTT(k)$ : **average round-trip time** for the first  $k$  segments

## ■ Expression

$$ARTT(k+1) = \frac{1}{k+1} \sum_{i=1}^{k+1} RTT(i) \quad \text{or}$$

$$ARTT(k+1) = \frac{1}{k+1} (k \times ARTT(k) + RTT(i))$$

$$\frac{k}{k+1} ARTT(k) + \frac{1}{k+1} RTT(i)$$



# Exponential Average


- Term
  - $SRTT(k)$ : **smoothed round-trip time** estimate for the first  $k$  segments
- Expression

$$SRTT(k+1) = \alpha \times SRTT(k) + (1-\alpha) \times RTT(k+1) \quad \text{i.e.}$$


$$SRTT(k+1) = (1-\alpha) \times RTT(k+1) + \alpha(1-\alpha) \times RTT(k) + \\ \alpha^2(1-\alpha) \times RTT(k-1) + \dots + \alpha^k(1-\alpha) \times RTT(1)$$



# Simple and Exponential Averaging



(a) Increasing function



(b) Decreasing function



# RFC 793

## ■ Term

- $RTO(k)$ : **retransmission timeout**, i.e. the timer after the first  $k$  segments

## ■ Expression

$$RTO(k+1) = \text{Min}(\text{UBOUND}, \text{MAX}(\text{LBOUND}, \beta \times SRTT(k+1)))$$

- Retransmission timer set between  $\text{LBOUND} \sim \text{UBOUND}$
- Suggested values,  $\alpha$ : 0.8~0.9,  $\beta$ : 1.3~2.0



# Jacobson's Algorithm (1)

## Problem in RFC 793

- Not counting **variance of RTT** (network stability)
- When network is stable, RTT variance is low, but  $\beta=1.3$  gives a higher RTO
- When network is **unstable**, RTT variance is high,  $\beta=2$  is inadequate to protect against retransmissions



## Jacobson's Algorithm (2)

### ■ Term

- $SERR(k)$ : **smoothed error estimate**, difference of round-trip time of segment  $k$  and the current  $SRTT$
- $SDEV(k)$ : **standard deviation** for round-trip time of first  $k$  segments

### ■ Expression

$$SRTT(k+1) = (1-g) \times SRTT(k) + g \times RTT(k+1)$$


$$SERR(k+1) = RTT(k+1) - SRTT(k)$$

$$SDEV(k+1) = (1-h) \times SDEV(k) + h \times |SERR(k+1)|$$


$$RTO(k+1) = SRTT(k+1) + f \times SDEV(k+1)$$

$$g = 1/8 = 0.125 \quad h = 1/4 = 0.25 \quad f = 2 \text{ or } 4$$

# Jacobson's RTO Calculation



(a) Increasing function



(b) Decreasing function





# Exponential RTO Backoff

- Timeout is often **due to congestion** by dropped packet or long round trip
  - Should slow down end system transmission
  - Maintaining RTO is not a good idea
- Similar to **Binary exponential backoff** in Ethernet
  - RTO multiplied each time a segment is re-transmitted
  - $RTO = q \times RTO$
  - Commonly  $q=2$



# Karn's Algorithm

- The problem
  - If a segment is **re-transmitted**, the *ACK* arriving may be
    - For the first copy of the segment, or for the second copy, or others
    - No way to tell
- RTT Sampling
  - Do **not measure RTT** for re-transmitted segments
  - Calculate RTO backoff when re-transmission occurs
  - Until *ACK* arrives for segment that has not been re-transmitted
  - Begin sampling, stop RTO backoff



# Window Management

- Add **congestion (send) window** besides the credit (receive window)

$$awnd = \text{Min}(\text{credit}, cwnd)$$

- **awnd**: allowed window, in **MSS (maximum segment size)**
- **credit**: the amount of unused credit granted in last **ACK**, in **MSS**
- **cwnd**: congestion window, in **MSS**

$$\text{Sending rate} = awnd \times MSS / RTT$$

- Manage congestion window
  - **Slow Start**: exponentially expending the **cwnd** at start of connection
  - **Dynamic Window on Congestion**: shrinking / expending the **cwnd** with stages when retransmission occurs



# TCP Reno

- Congestion Window Management for TCP
- Proposed by Jacobson, 1990
- Three stages:
  - (1) Slow Start
  - (2) Congestion Avoidance
  - (3) Fast Retransmit and Fast Recovery




# Slow Start

## (1) 慢启动



- When connection begins,  $cwnd = 1 MSS$
- Each time an *ACK* received,  $cwnd$  increased by ACKed number of *MSSs* until Max value reached
- $cwnd$  increased exponentially until first loss event occurs
  - Timeout or 3 duplicate *ACKs*



Summary: initial rate is slow but ramps up exponentially fast



# Dynamic Windows Sizing

- By Jacobson, set slow start threshold **ssthresh** =  $cwnd/2$
- After 3 duplicate **ACKs** (丢包事件)
  - Network still capable of delivering some segments
  - $cwnd$  is set to be **ssthresh**
  - Enter congestion avoidance:  $cwnd$  increases by 1 (linearly instead of exponentially) after each RTT or **ACK** received
- After **timeout event** (超时事件)
  - $cwnd$  is set to 1
  - Back to slow start again
- **Else:** If  $cwnd$  reaches **ssthresh**
  - Enter congestion avoidance:  $cwnd$  increases by 1 after each RTT or **ACK** received




# Congestion Avoidance

(2) 拥塞避免及快恢复:  
加性增, 乘性减 (AIMD)

- ❖ **Approach:** sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
  - **additive increase:** increase **cwnd** by 1 MSS every RTT until loss detected
  - **multiplicative decrease:** cut **cwnd** in half after loss ( 3 duplicate ACKs )

AIMD saw tooth behavior: probing for bandwidth





# Illustration of Window Management




Figure 17.14 Illustration of Slow Start and Congestion Avoidance






# Fast Retransmit and Fast Recovery

## ■ Fast Retransmit

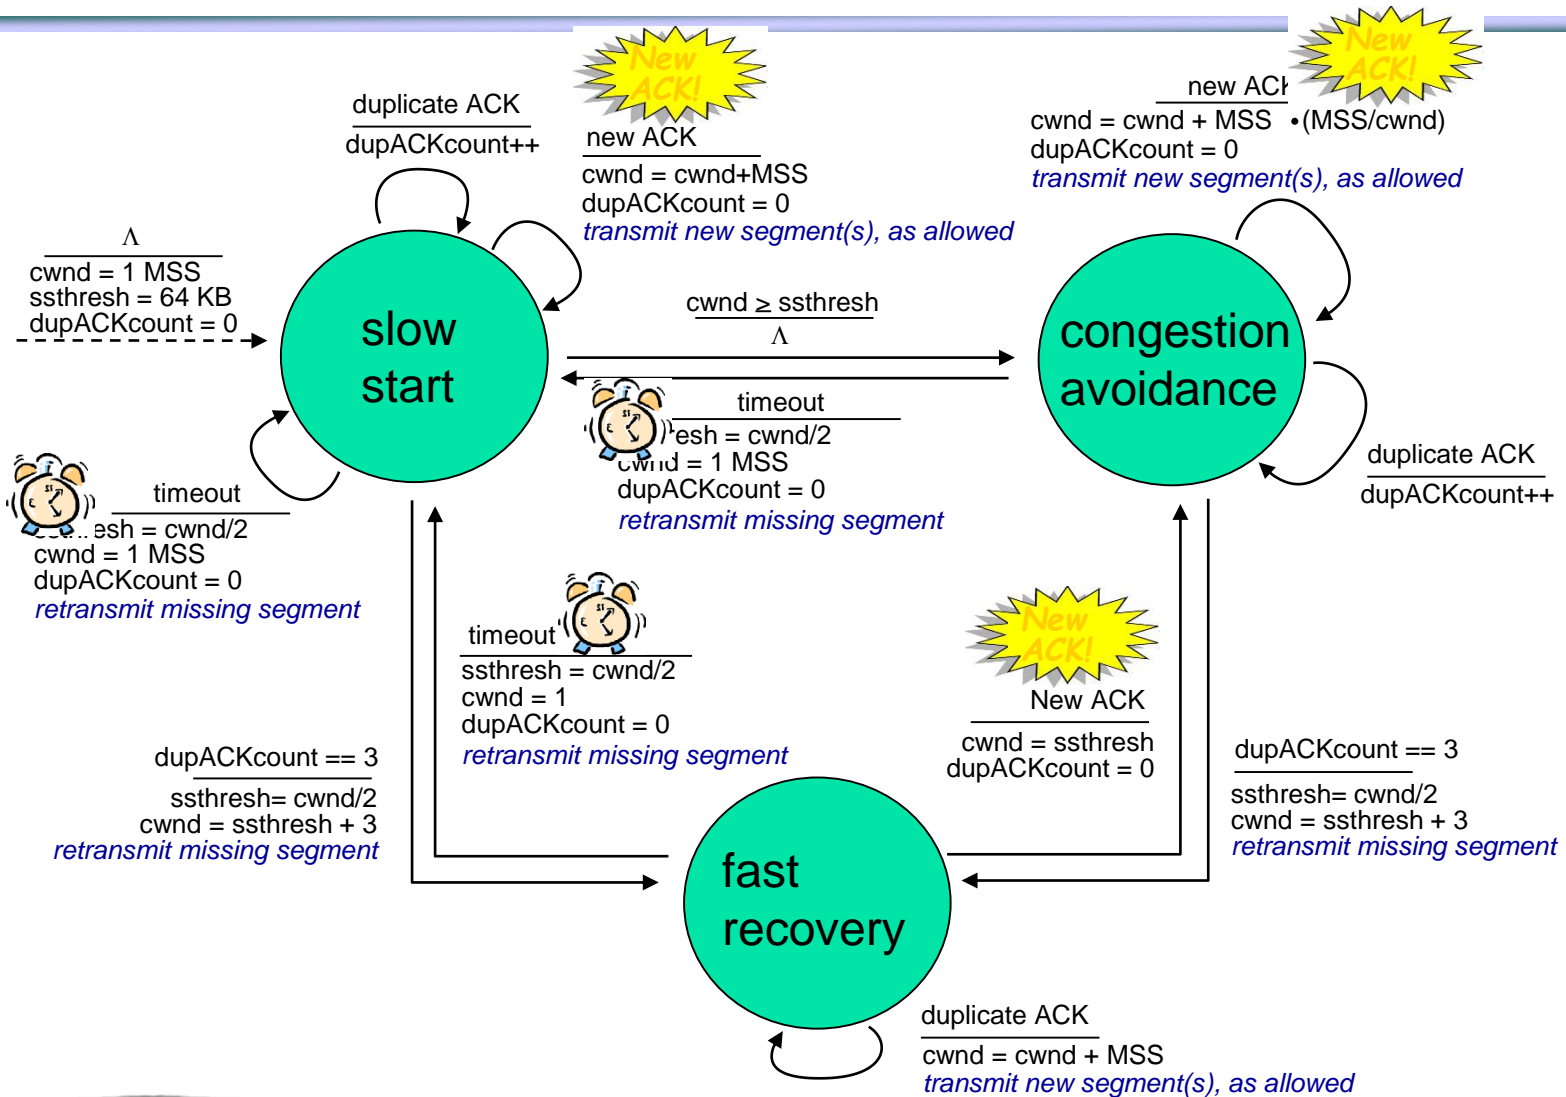
- If sender receives 3 ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #
  - likely that unacked segment lost, so don’t wait for timeout

## ■ Fast Recovery

- If sender receives 3 ACKs for same data, set  $ssthresh = ssthresh / 2$
- (TCP Reno:) Set  $cwnd = cwnd / 2$ , enter congestion avoidance:  $cwnd$  increases by 1 after each RTT or ACK received. (TCP Tahoe will set  $cwnd = 1$  and enters slow start)



# Summary: TCP Congestion Control





# History

- Original TCP
  - Only flow control, no congestion control
  - Only consider receiver's capacity, not consider network capacity
  - Cause congestion collapse (1986)
- 1988, Jacobson, TCP Tahoe
  - Slow start, AIMD, Fast retransmit
- 1990, Jacobson, TCP Reno
  - Slow start, AIMD, Fast retransmit+Fast recovery
- 1990-, new variations
  - TCP-NewReno
  - TCP SACK
  - TCP Vegas




# New Window Management Algorithm

- Problem of **linear increase**
  - Many **long fat networks**: large bandwidth with long delay
  - Slow response of TCP in such networks leaves sizeable unused bandwidth
- An example
  - A TCP connection with *1250-Octet MSS* (*Maxitum Segment Size* ) and *100ms RTT* on *10Gbps* network
  - To fully use the network, credit is big, and nearly **1.4 hour** is needed for linear increase



# Too Slow Linear Increase

- Q: how to improve it





# BIC and CUBIC

- **BIC** (Binary Increase Congestion control)
  - Implemented and used by default in Linux kernels 2.6.8
- **CUBIC**
  - The window is a cubic function of time since the last congestion event
  - Implemented and used by default in Linux kernels 2.6.19 and above



- BIC adaptively increase cwnd, and decrease cwnd by 1/8

$$cwnd = cwnd + 1$$




$$cwnd = cwnd + f(cwnd, history)$$

$$cwnd = cwnd * (1 - 1/2)$$



$$cwnd = cwnd * (1 - 1/8)$$





# BIC Overview

- 2 stages
  - **Binary Search**: increase window after congestion
  - **Max Probing**: search for better window size (until max credit)
    - (Max window size if set from history information. If there are more size available, Max\_Probing is used for the purpose of exploring larger window size)
- 4 parameters defined
  - ***S<sub>max</sub>*** : the maximum increment, e.g.  $1/8 \times credit$
  - ***S<sub>min</sub>*** : the minimum increment, e.g.  $2 \text{ MSS}$ 
    - (Using  $S_{max}$ ,  $S_{min}$  to avoid jitter: If a binary search steps too large, the traffic will change quickly. Thus  $S_{max}$  is used to limit the maximum change in one step)
  - ***W<sub>max</sub>*** : maximum window size of current search, e.g. window size just before the lost
  - ***W<sub>min</sub>*** : minimum window size of current search, e.g. current window size without lost





# BIC Stages

## ■ Additive Increase

- Linear increase with  $inc$


## ■ Binary search

- How to set  $inc$

```
while ( $Wmin \leq Wmax$ ){  
     $inc = (Wmin + Wmax) / 2 - cwnd$ ;  
    if ( $inc > Smax$ )  $inc = Smax$ ;  
    else if ( $inc < Smin$ )  $inc = Smin$ ;  
     $cwnd += inc$ ;  
  
    if (no packet losses)  $Wmin = cwnd$ ;  
    else {  
         $Wmax = cwnd$ ;  
         $Wmin = cwnd \times \beta$  (e.g. 0.8)  
    }  
}
```

# BIC Stages

- **Binary Search Stage:** Additive increase + Binary search
- **Max Probing Stage:** Binary search + Additive increase



At the beginning, each step increases  $S_{\max}$

In the middle, pure binary search

To probing more window size, each step increases by binary search or  $S_{\min}$



# CUBIC

- BIC problem
  - The BIC's growth function may be **too aggressive** for TCP
  - BIC is not suitable for short RTT or low speed networks (may cause unfairness)
- Handle
  - Express the multi-stage BIC curve with a **single cubic function**



# CUBIC Overview

## ■ Parameters


- ***W<sub>cubic</sub>***: current cwnd
- ***W<sub>max</sub>***: window size just before the last lost
- ***T***: elapsed time from the last lost
- ***C***: a scaling constant

■ Function  $W_{cubic} = C(T - K)^3 + W_{max}$

$$\text{where } K = \sqrt[3]{\frac{W_{max} \times (1 - \beta)}{C}}$$



# CUBIC Overview





# Summary

- UDP & TCP
- TCP header fields
- TCP congestion control
  - Retransmission timer
  - Window management



# Homework

- 第三章: R14, P27, P32, P40, P45, P46, P50, P52